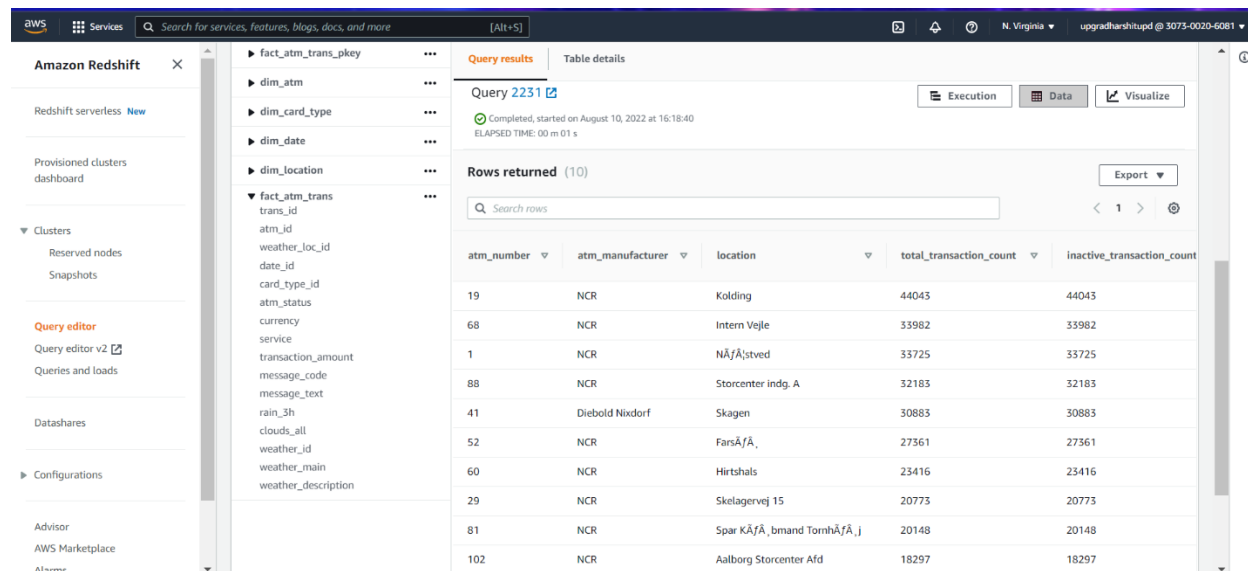


Solving analytical queries on Redshift Cluster

Here, you have to write the query used for solving the question and the screenshots of the table which is outputted after the query is run on the AWS Redshift Query editor UI.

1. Top 10 ATMs where most transactions are in the 'inactive' state

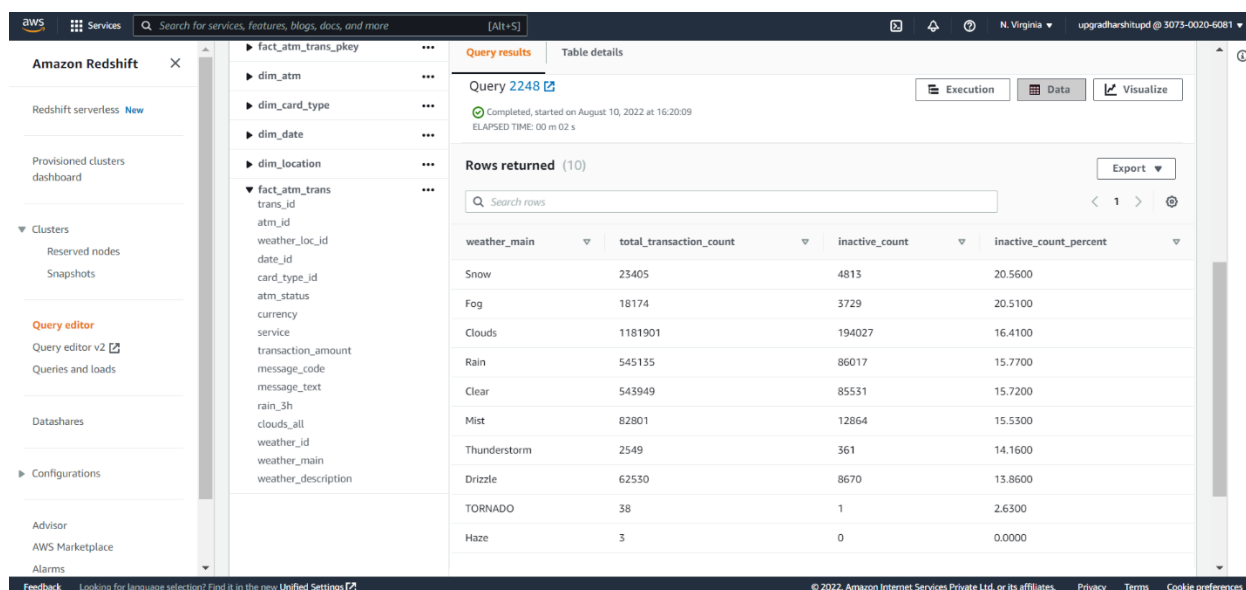
```
select a.atm_number, a.atm_manufacturer, l.location,
count(trans_id) as total_transaction_count,
sum(case when atm_status = 'Inactive' then 1 else 0 end) as inactive_transaction_count,
(inactive_transaction_count/total_transaction_count)*100 as count_percent
from atm_data.fact_atm_trans f, atm_data.dim_atm a, atm_data.dim_location l
where f.atm_id = a.atm_id and a.atm_location_id = l.location_id
group by a.atm_number, a.atm_manufacturer, l.location
having count_percent > 50
order by inactive_transaction_count desc
limit 10;
```



atm_number	atm_manufacturer	location	total_transaction_count	inactive_transaction_count
19	NCR	Kolding	44043	44043
68	NCR	Intern Vejle	33982	33982
1	NCR	NÅfÅstved	33725	33725
88	NCR	Storcenter indg. A	32183	32183
41	Diebold Nixdorf	Skagen	30883	30883
52	NCR	FarsÅfÅ	27361	27361
60	NCR	Hirtshals	23416	23416
29	NCR	Skelagervej 15	20773	20773
81	NCR	Spar KÅfÅ, bmand TornhÅfÅ, j	20148	20148
102	NCR	Aalborg Storcenter Afd	18297	18297

2. Number of ATM failures corresponding to the different weather conditions recorded at the time of the transactions

```
select f.weather_main,
count(trans_id) as total_transaction_count,
sum(case when atm_status = 'Inactive' then 1 else 0 end) as inactive_count,
case when coalesce(inactive_count, 0) = 0 then 0.0000
else trunc((cast(inactive_count as numeric(10,4))/total_transaction_count)*100, 2)
end as inactive_count_percent
from atm_data.fact_atm_trans f
where f.weather_main != "
group by f.weather_main
order by inactive_count_percent desc
limit 10;
```



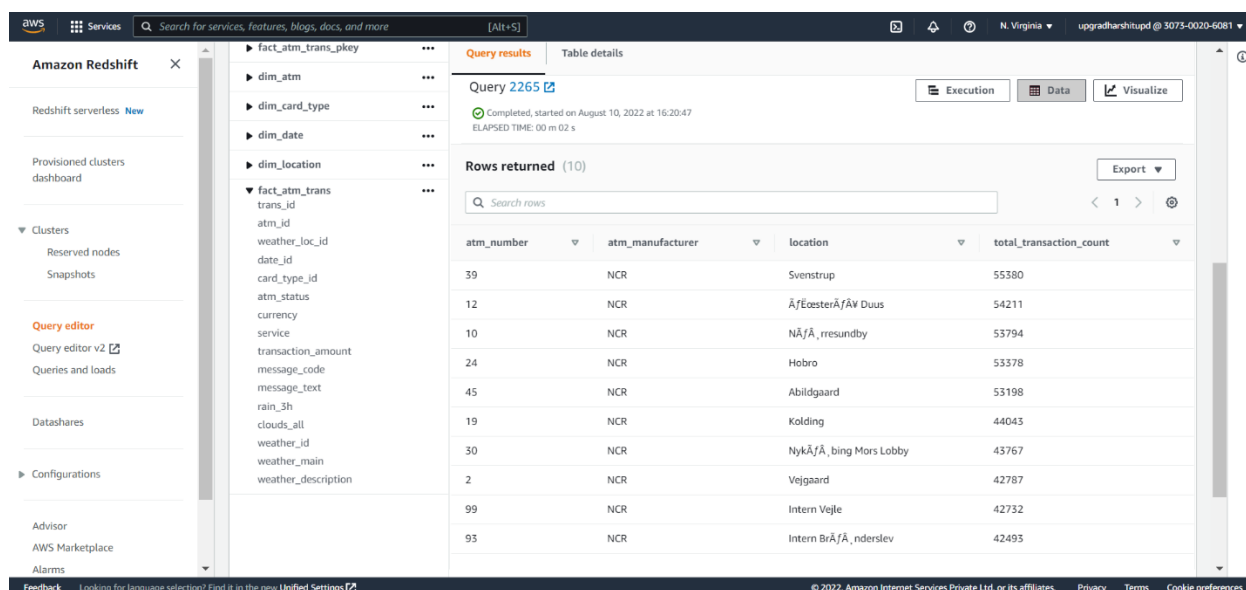
Query 2248 [Link]
Completed, started on August 10, 2022 at 16:20:09
ELAPSED TIME: 00 m 02 s

Rows returned (10)

weather_main	total_transaction_count	inactive_count	inactive_count_percent
Snow	23405	4813	20.5600
Fog	18174	3729	20.5100
Clouds	1181901	194027	16.4100
Rain	545135	86017	15.7700
Clear	543949	85531	15.7200
Mist	82801	12864	15.5300
Thunderstorm	2549	361	14.1600
Drizzle	62530	8670	13.8600
TORNADO	38	1	2.6300
Haze	3	0	0.0000

3. Top 10 ATMs with the most number of transactions throughout the year

```
select a.atm_number, a.atm_manufacturer, l.location,
count(trans_id) as total_transaction_count
from atm_data.fact_atm_trans f, atm_data.dim_atm a, atm_data.dim_location l
where f.atm_id = a.atm_id and a.atm_location_id = l.location_id
group by a.atm_number, a.atm_manufacturer, l.location
order by total_transaction_count desc
limit 10;
```

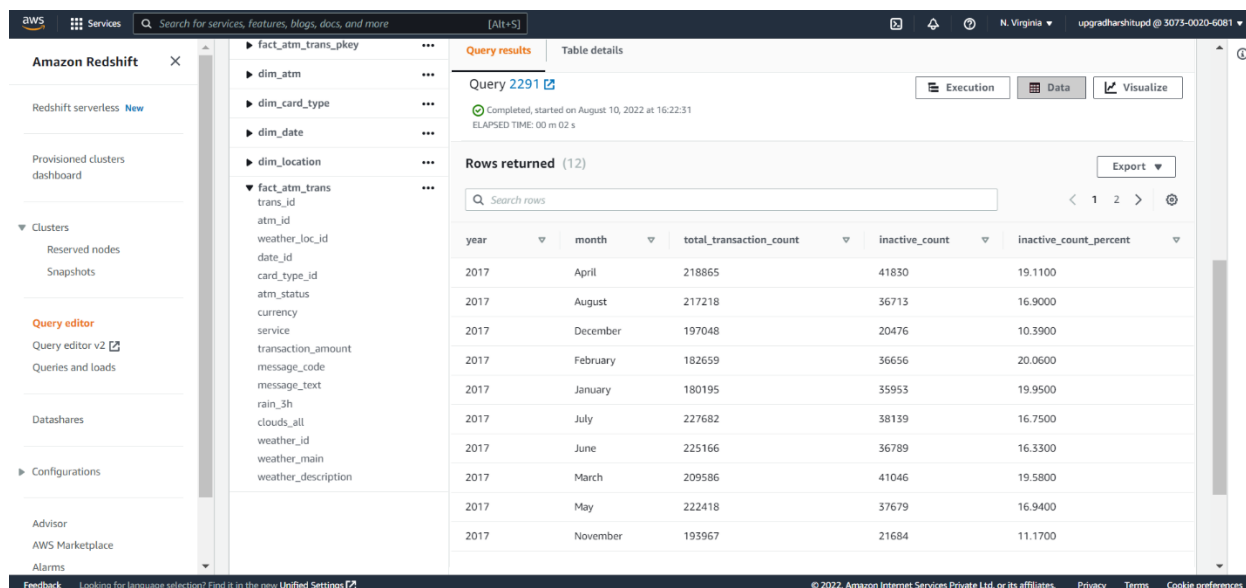


The screenshot shows the AWS Redshift Query Results interface. The query results are displayed in a table with the following columns: atm_number, atm_manufacturer, location, and total_transaction_count. The results are sorted by total_transaction_count in descending order.

atm_number	atm_manufacturer	location	total_transaction_count
39	NCR	Svenstrup	55380
12	NCR	Århus	54211
10	NCR	Næstved	53794
24	NCR	Hobro	53378
45	NCR	Abildgaard	53198
19	NCR	Kolding	44043
30	NCR	Nykøbing Mors	43767
2	NCR	Vejgaard	42787
99	NCR	Intern Vejle	42732
93	NCR	Intern Brønderslev	42493

4. Number of overall ATM transactions going inactive per month for each month

```
select d.year, d.month,
count(trans_id) as total_transaction_count,
sum(case when atm_status = 'Inactive' then 1 else 0 end) as inactive_count,
case when coalesce(inactive_count, 0) = 0 then 0.0000
else trunc((cast(inactive_count as
numeric(10,4))/total_transaction_count)*100, 2)
end as inactive_count_percent from atm_data.fact_atm_trans f inner join atm_data.dim_date d
on f.date_id = d.date_id
group by d.year, d.month
order by d.year, d.month
```

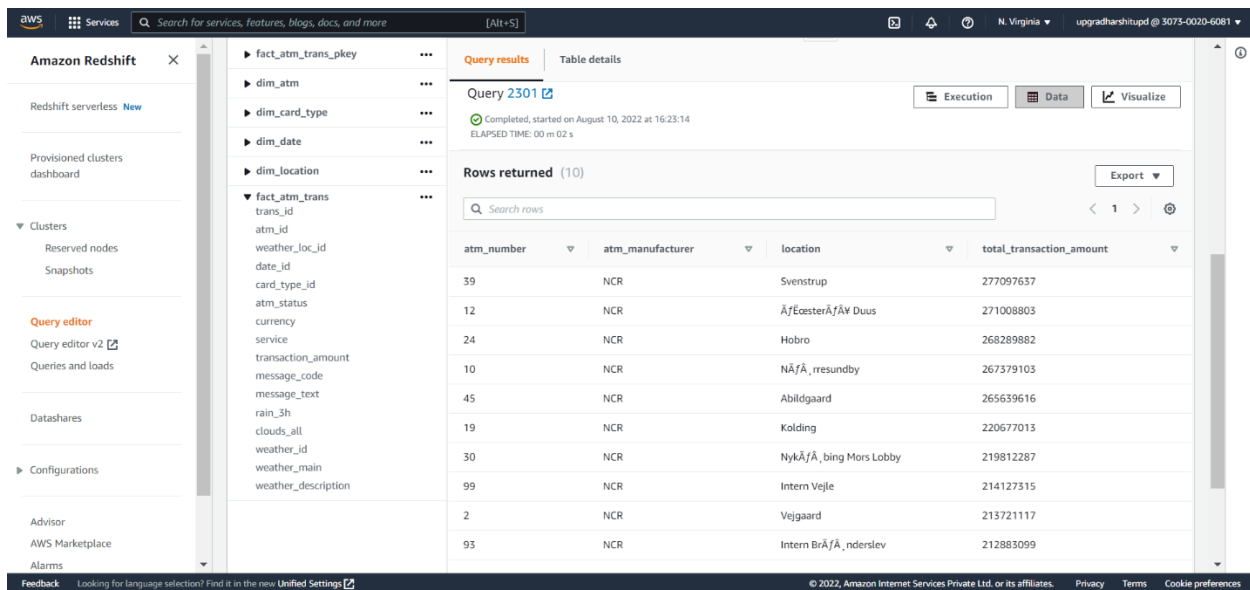


Query results for Query 2291. The query is completed and returned 12 rows of data.

year	month	total_transaction_count	inactive_count	inactive_count_percent
2017	April	218865	41830	19.1100
2017	August	217218	36713	16.9000
2017	December	197048	20476	10.3900
2017	February	182659	36656	20.0600
2017	January	180195	35953	19.9500
2017	July	227682	38139	16.7500
2017	June	225166	36789	16.3300
2017	March	209586	41046	19.5800
2017	May	222418	37679	16.9400
2017	November	193967	21684	11.1700

5. Top 10 ATMs with the highest total withdrawn amount throughout the year

```
select a.atm_number, a.atm_manufacturer, l.location,
sum(transaction_amount) as total_transaction_amount
from atm_data.fact_atm_trans f, atm_data.dim_atm a, atm_data.dim_location l
where f.atm_id = a.atm_id and a.atm_location_id = l.location_id
group by a.atm_number, a.atm_manufacturer, l.location
order by total_transaction_amount desc
limit 10;
```

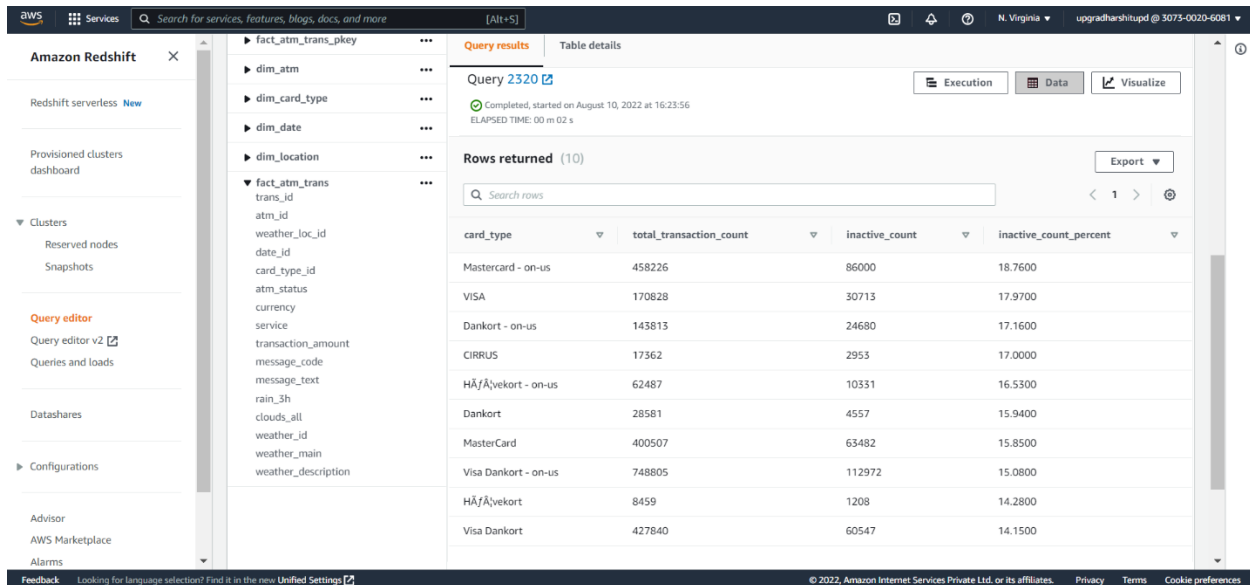


The screenshot shows the Amazon Redshift Query Results interface. The query is titled "Query 2301" and was completed on August 10, 2022, at 16:23:14. The elapsed time is 00 m 02 s. The results show 10 rows returned, sorted by total_transaction_amount in descending order.

atm_number	atm_manufacturer	location	total_transaction_amount
39	NCR	Svenstrup	277097637
12	NCR	ÅrEøsterÅr Duus	271008803
24	NCR	Hobro	268289882
10	NCR	NÅrÅ, resundby	267379103
45	NCR	Abildgaard	265639616
19	NCR	Kolding	220677013
30	NCR	NykÅrÅ, bing Mors Lobby	219812287
99	NCR	Intern Vejle	214127315
2	NCR	Vejgaard	213721117
93	NCR	Intern BrÅrÅ, nderslev	212883099

6. Number of failed ATM transactions across various card types

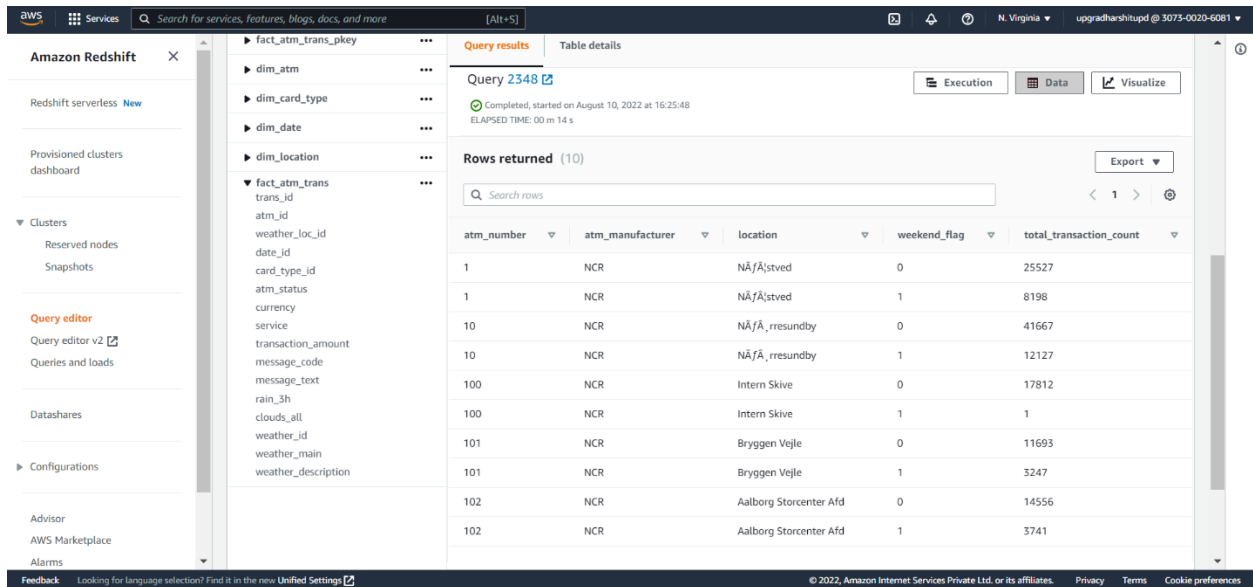
```
select ct.card_type,
count(trans_id) as total_transaction_count,
sum(case when atm_status = 'Inactive' then 1 else 0 end) as inactive_count,
case when coalesce(inactive_count, 0) = 0 then 0.0000
else trunc((cast(inactive_count as numeric(10,4))/total_transaction_count)*100, 2)
end as inactive_count_percent
from atm_data.fact_atm_trans f, atm_data.dim_card_type ct
where f.card_type_id = ct.card_type_id
group by ct.card_type
order by inactive_count_percent desc
limit 10;
```



card_type	total_transaction_count	inactive_count	inactive_count_percent
Mastercard - on-us	458226	86000	18.7600
VISA	170828	30713	17.9700
Dankort - on-us	143813	24680	17.1600
CIRRUS	17362	2953	17.0000
HÃ¶fÃ¶vekort - on-us	62487	10331	16.5300
Dankort	28581	4557	15.9400
MasterCard	400507	63482	15.8500
Visa Dankort - on-us	748805	112972	15.0800
HÃ¶fÃ¶vekort	8459	1208	14.2800
Visa Dankort	427840	60547	14.1500

7. Number of transactions happening on an ATM on weekdays and on weekends throughout the year. Order this by the ATM_number, ATM_manufacturer, location, weekend_flag and then total_transaction_count

```
select a.atm_number, a.atm_manufacturer, l.location,
case when d.weekday in ('Saturday','Sunday') then 1 else 0 end as weekend_flag,
count(trans_id) as total_transaction_count
from atm_data.fact_atm_trans f, atm_data.dim_atm a, atm_data.dim_location l,
atm_data.dim_date d
where f.atm_id = a.atm_id and a.atm_location_id = l.location_id and f.date_id = d.date_id
group by a.atm_number, a.atm_manufacturer, l.location, weekend_flag
order by a.atm_number, a.atm_manufacturer, l.location, weekend_flag, total_transaction_count
limit 10;
```



Query results

Query 2348

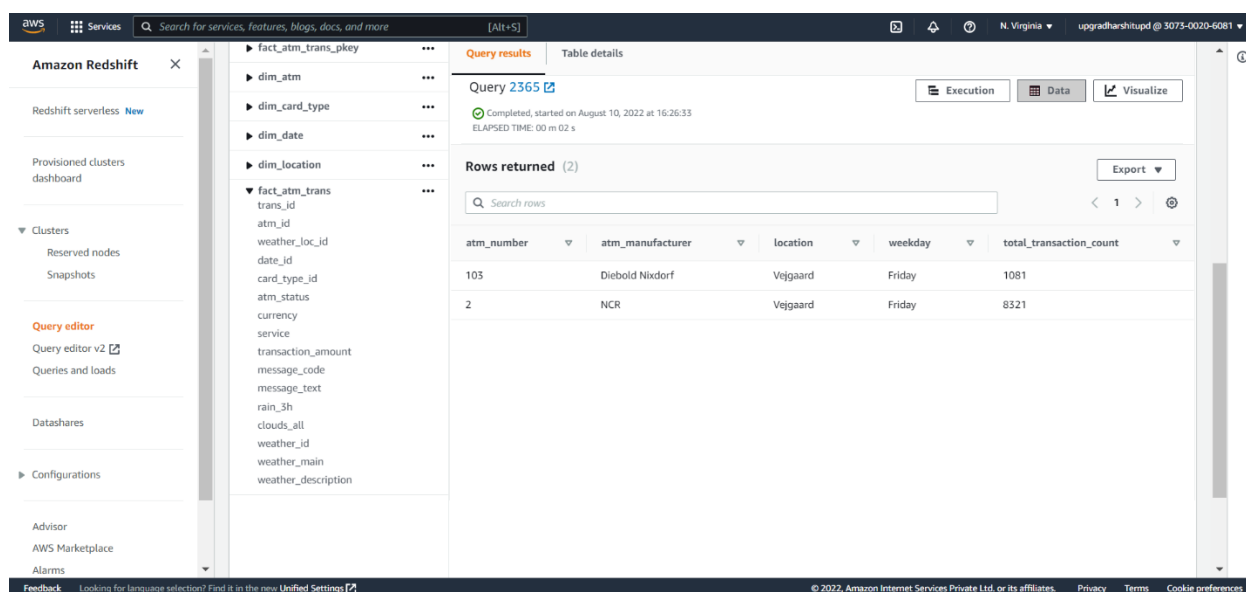
Completed, started on August 10, 2022 at 16:25:48
ELAPSED TIME: 00 m 14 s

Rows returned (10)

atm_number	atm_manufacturer	location	weekend_flag	total_transaction_count
1	NCR	NÅfÅstved	0	25527
1	NCR	NÅfÅstved	1	8198
10	NCR	NÅfÅ ,resundby	0	41667
10	NCR	NÅfÅ ,resundby	1	12127
100	NCR	Intern Skive	0	17812
100	NCR	Intern Skive	1	1
101	NCR	Bryggen Vejle	0	11693
101	NCR	Bryggen Vejle	1	3247
102	NCR	Aalborg Storcenter Afd	0	14556
102	NCR	Aalborg Storcenter Afd	1	3741

8. Most active day in each ATMs from location "Vejgaard"

```
select a.atm_number, a.atm_manufacturer, l.location, d.weekday,
count(trans_id) as total_transaction_count
from atm_data.fact_atm_trans f inner join atm_data.dim_atm a on f.atm_id = a.atm_id
inner join atm_data.dim_location l on a.atm_location_id = l.location_id
inner join atm_data.dim_date d on f.date_id = d.date_id
where l.location = 'Vejgaard' and d.weekday in
( select d.weekday
from atm_data.fact_atm_trans f inner join atm_data.dim_date d
on f.date_id = d.date_id
inner join atm_data.dim_location l on f.weather_loc_id = l.location_id
where l.location = 'Vejgaard'
group by d.weekday
order by count(f.trans_id) desc
limit 1 )
group by a.atm_number, a.atm_manufacturer, l.location, d.weekday
order by total_transaction_count;
```



The screenshot shows the AWS Redshift Query Editor interface. The left sidebar contains navigation options like 'Amazon Redshift', 'Redshift serverless', 'Provisioned clusters', 'Clusters', 'Query editor', 'Databases', 'Configurations', 'Advisor', 'AWS Marketplace', and 'Alarms'. The main panel displays the 'Query results' for 'Query 2365', which is 'Completed, started on August 10, 2022 at 16:26:33' with an 'ELAPSED TIME: 00 m 02 s'. Below the query status, it shows 'Rows returned (2)' and a table with the following data:

atm_number	atm_manufacturer	location	weekday	total_transaction_count
103	Diebold Nixdorf	Vejgaard	Friday	1081
2	NCR	Vejgaard	Friday	8321